

Technical Overview

May 25, 2023

Team name: Storskov inc.

Company: Systematic

Teaching assistant: Mikkel Lippert

Agile coach: Chamara

Role	email	Name
PO	Miboj@systematic.com	Mico Bøje
SM	isam@itu.dk	Isabella Amstrup
dev	ahad@itu.dk	Adam Temsamani
dev	bath@itu.dk	Baldur Thomsen
dev	ehel@itu.dk	Emilia Helsted
dev	jakst@itu.dk	Jakob Storskov
dev	jklo@itu.dk	Johan Skoven
dev	kmsa@itu.dk	Katrine Sandø
dev	masc@itu.dk	Malte Schlack
dev	biha@itu.dk	Maxime Havez
dev	olfw@itu.dk	Oliver Wilhjelm

Contents

1	Product	3
2	Software Qualities	3
2.1	Ensuring Software Qualities	3
3	Architecture Design	4
4	Technical Specifications	5
5	appendix	7
5.1	Usertest	7

1 Product

The product is an anomaly detection service, that uses a machine learning model to analyse given log-messages for anomalies. Users are then informed and given the option to handle them through a web-app.

2 Software Qualities

Below is a list of software qualities we intend to uphold throughout the project.

Usability:

The dashboard should be intuitive to use for IT professionals. Important information such as alerts and dashboard should be the first thing users are presented with.

Performance:

The system should be able to handle at least 3.000 log messages per minute.

Security:

Logins must be handled in a secure fashion, meaning that passwords are at least hashed.

Scalability:

The system should be able to handle increasing workloads by scaling components. This means that everything except the database should ideally be able to run in multiple instances.

Maintainability:

The code must be documented, commented and preferably achieve at least a B in maintainability on a Sonarqube scan.

Testability:

The code should be tested in a CI/CD pipeline, like GitHub workflow.

2.1 Ensuring Software Qualities

Usability:

To ensure good usability, we have done user testing¹ and used the given feedback to improve the user-experience

Performance:

The product has achieved this software quality (assuming adequate hardware) through use of the RabbitMQ framework's load-balancing abilities.

Security:

The product uses the Keycloak solution to handle user management and authorisation to achieve this. Further masking of sensitive information is considered a nice to have feature, which could be solved using Istio.

Scalability:

The product has achieved this software quality through use of the RabbitMQ

¹5.1

framework's load-balancing abilities.

Maintainability:

We have set up pre-commit hooks and GitHub actions to help with maintainability, along with a pull request template. Together, they enforce and remind us of upholding a clean codebase for easier maintainability. The codebase has been thoroughly documented and the final product has achieved acceptable results in a SonarQube scan.

Testability:

Testing has primarily been done through API-calls via Postman, as most of our program is dependent on these. The front-end has been tested further through use of user-tests.

3 Architecture Design

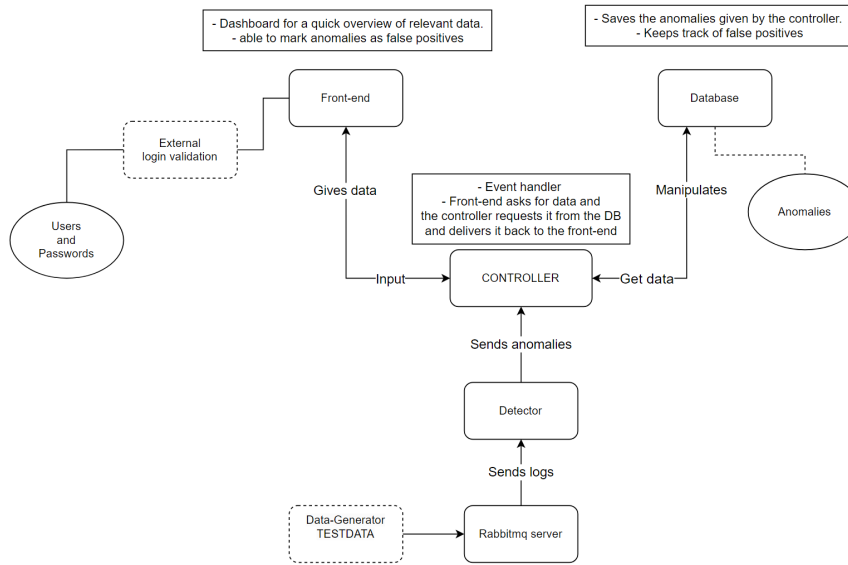


Figure 1: Overall architecture of the system, updated for exam submission

Figure 1 shows the software architecture of the program. We have based our architecture on a Model View Controller architecture.

The system uses a machine learning algorithm build in Python to analyse log messages. The model we got from our PO is written in Python. It was therefore decided that the primary language for the entire codebase is Python to maintain consistency.

The system uses a Postgres database, to store our data. This database was

given to us along with the initial codebase, and there was no reason to change the type of database. The system uses the open source solution Keycloak for user management, as Systematic is already using it for other products. It also uses RabbitMQ as a message broker. This ensures no messages are lost in case of a system crash, and enables load balancing by having multiple active anomaly detectors. The RabbitMQ framework can hence be used to scale the product horizontally. Lastly the system uses Docker for containerization as it is a industry standard and was desired by our PO.

4 Technical Specifications

The system is divided into a front-end and back-end. The front-end is on a single docker container and made primarily with python dash and pandas frameworks. This includes functions and callbacks with API-calls that connects the front-end to the back-end.

The back-end is further divided into 6 different docker containers, handling different parts of the system: anomaly-detector, controller, Keycloak, RabbitMQ, Postgres, and a data-generator.

Anomaly-detector

The anomaly-detector contains the machine-learning model which handles analyzing given log-messages for anomalies. If the analyzed message is considered an anomaly it is given to the controller through an API call.

Controller

The controller handles communication between the front-end and the back-end, with the exception of the Keycloak container. The controller also handles communication between the anomaly-detector and the database. This is done through API-calls, which are made with the FastAPI framework.

Keycloak

The Keycloak container handles user-management authentication of user credentials. The container uses the open-source Keycloak solution, which is made for handling user identity and access. The Keycloak container communicates directly with the front-end through methods made using the Python Keycloak library.

RabbitMQ

The RabbitMQ container handles message-queuing and load-balancing through the use of the open source RabbitMQ message broker software. The container

divides the data between the amount of initialized anomaly-detectors. It also prevents data loss in case of system failure.

Postgres database

The Postgres container is our Postgres database, which stores anomaly-messages including all relevant information. It is accessed by the controller through CRUD operations.

Data-generator

The data-generator sends a constant stream of data to the RabbitMQ container. It currently simulates real life data but it is intended to be replaced actual data eventually.

5 appendix

5.1 Usertest

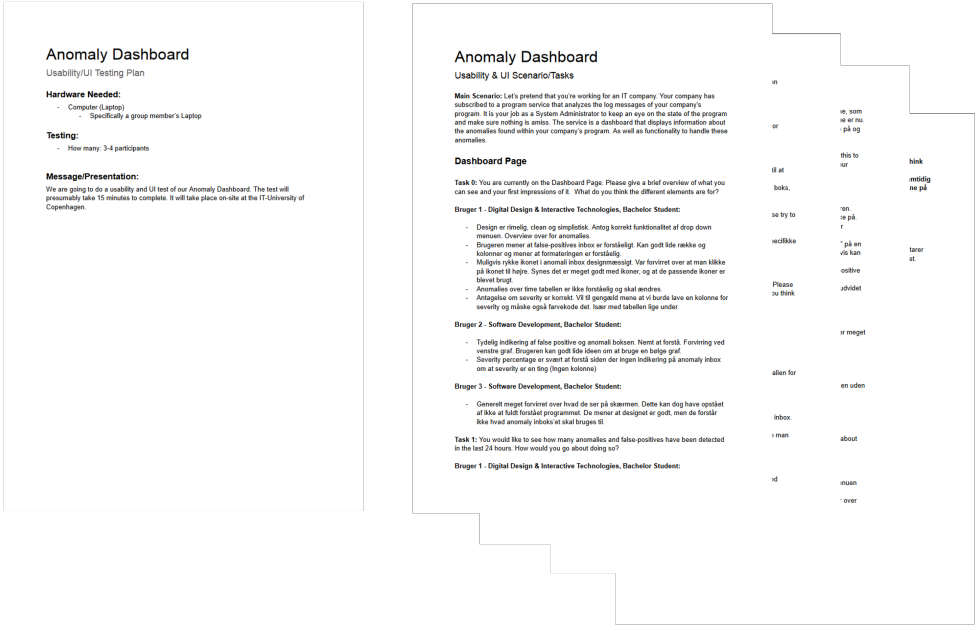


Figure 2: Questions for user tests

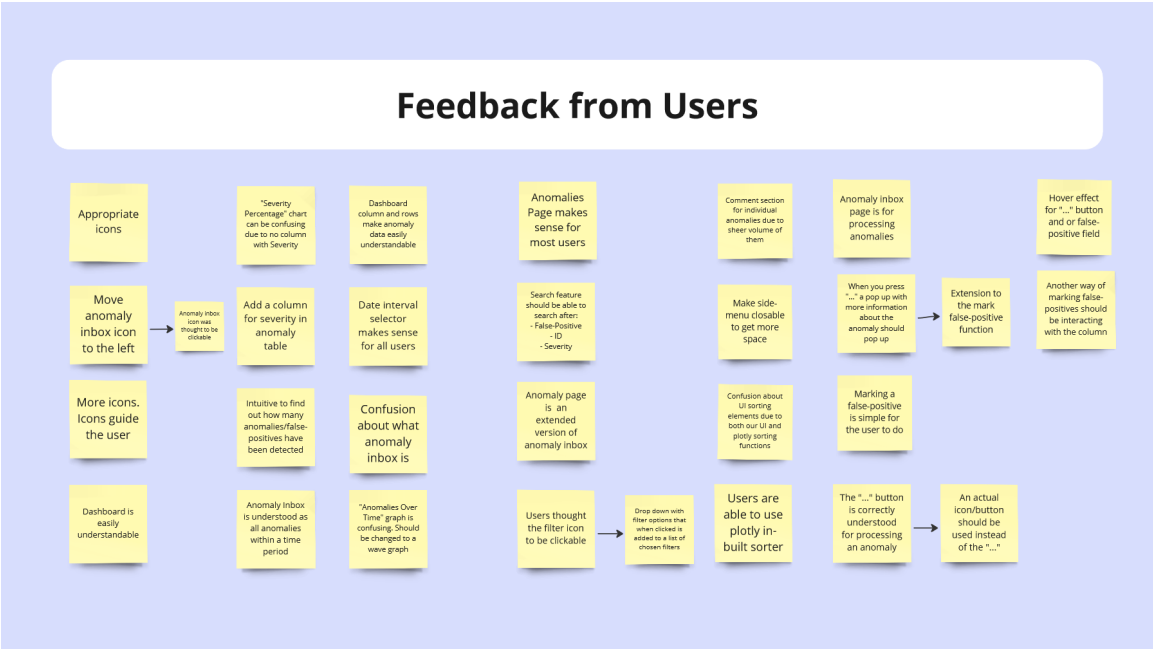


Figure 3: Feedback from users

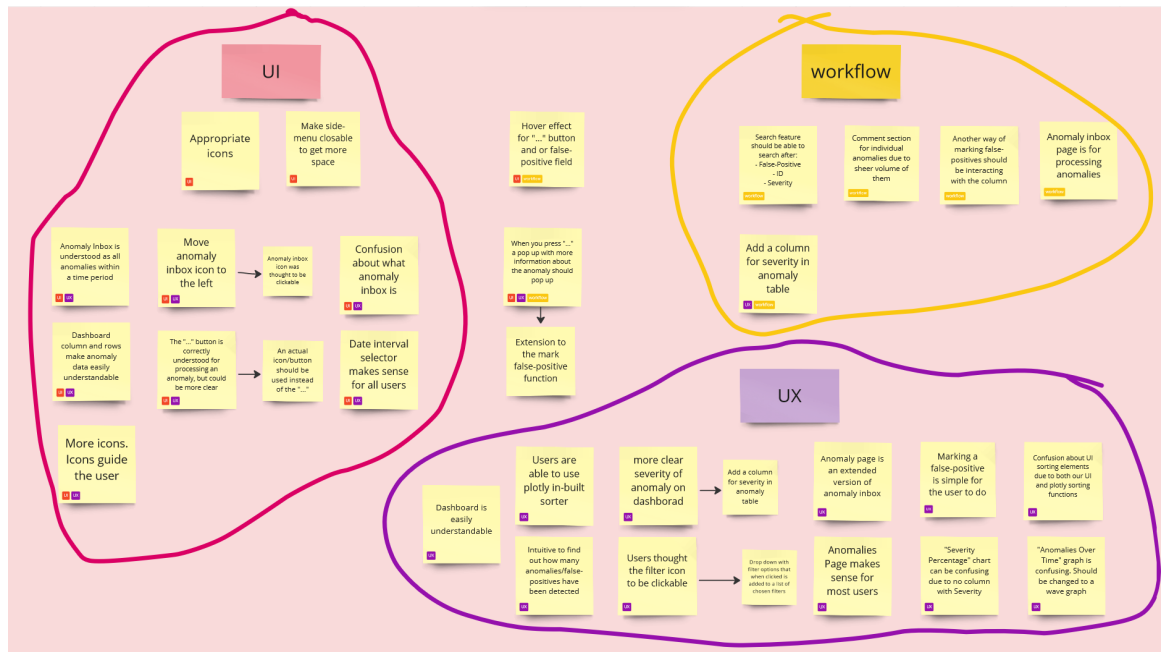


Figure 4: Grouping of user data

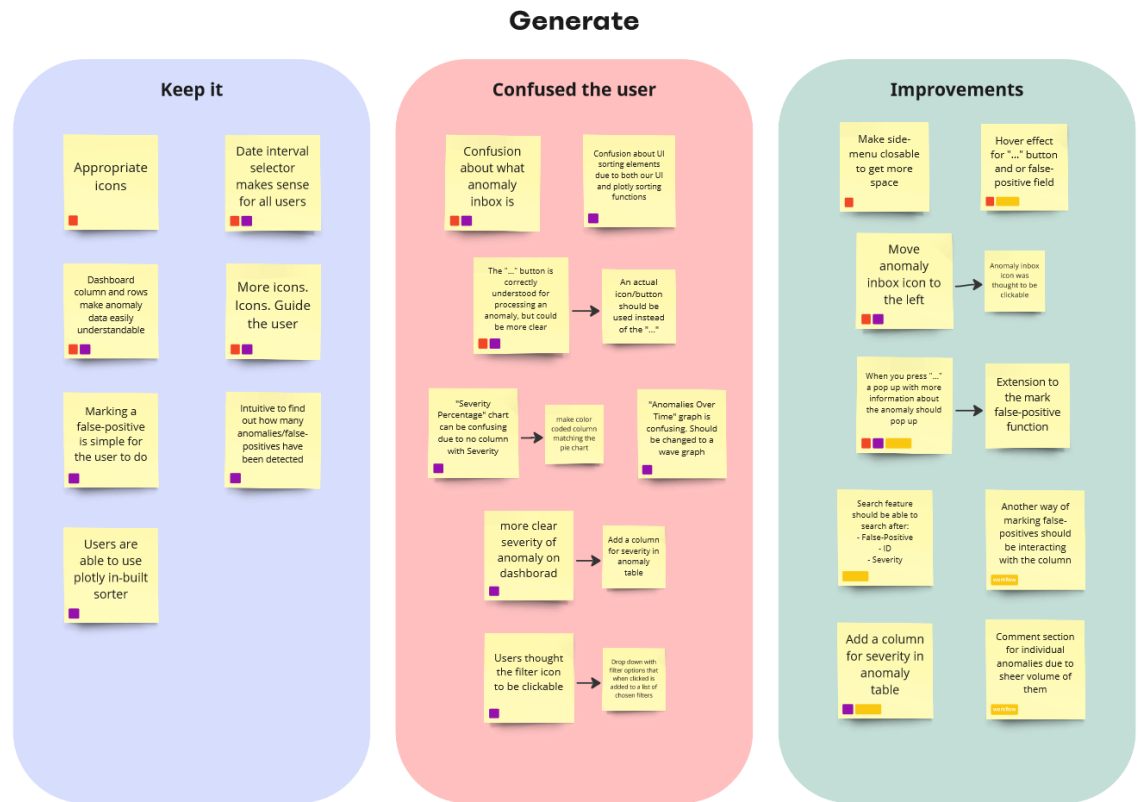


Figure 5: Final results of processing the user feedback from the user tests.

References

- [1] Kent Beck, Mike Beedle, Arie van Bennekum, Alistair Cockburn, Ward Cunningham, Martin Fowler, James Grenning, Jim Highsmith, Andrew Hunt, Ron Jeffries, Jon Kern, Brian Marick, Robert C. Martin, Steve Mellor, Ken Schwaber, Jeff Sutherland, and Dave Thomas. Manifesto for agile software development, 2001.